



# **Progress reporting in Postgres**

Álvaro Herrera  
PostgreSQL developer  
PgConf.Brasil, July 2019

# Progress Reporting

- Reporting of what?
- How does it look?
- How to use it?
- What commands are supported?
- How can I implement more?

# Progress Reporting

- Reporting of what?
- How does it look?
- How to use it?
- What commands are supported?
- How can I implement more?



# DDL progress reporting

- Many DDL commands take very long time to execute
  - VACUUM, CREATE INDEX, etc
- It's useful to have insight as to:
  - How much total work there is
  - How much work we have already done
- Allows to extrapolate

# DDL progress reporting

- Many DDL commands take very long time to execute
  - VACUUM, CREATE INDEX, etc
- It's useful to have insight as to:
  - How much total work there is
  - How much work we have already done
- Allows to extrapolate
- ... with caveats




*5% Complete*

# Feature design principles

- We want to present hard facts
- Not fiction
  - No guessing
  - No busted percentages
    - 0% – 95% in one minute ... *then a slow crawl to 99%*
    - ... *245% done*
    - *progress bars going backwards*

# Feature design principles

- We want to present hard facts
- Not fiction
  - No guessing
  - No busted percentages
    - 0% – 95% in one minute ... *then a slow crawl to 99%*
    - ... *245% done*
    - *progress bars going backwards*
- ... Preferably, detailed and useful facts



# Reporting VACUUM progress

- PostgreSQL 10
- Add a generic command progress reporting facility.
  - <http://git.postgresql.org/pg/commitdiff/b6fb6471f6af>
  - Vinayak Pokale, Rahila Syed, Amit Langote, Robert Haas.
- Add simple VACUUM progress reporting.
  - <http://git.postgresql.org/pg/commitdiff/c16dc1aca5e0>
  - Amit Langote, Robert Haas, Vinayak Pokale, Rahila Syed.



# Reporting vacuum progress (2)

```
alvherre=# SELECT * FROM pg_stat_progress_vacuum;
```

*Record 1*

pid	4204
datid	12386
datname	alvherre
relid	234754
phase	scanning heap
heap_blks_total	89759
heap_blks_scanned	61181
heap_blks_vacuumed	0
index_vacuum_count	0
max_dead_tuples	291
num_dead_tuples	0

# Reporting vacuum progress (3)

```
alvherre=# SELECT now(), pid, relid::regclass as table, phase,  
heap_blks_total, heap_blks_scanned, heap_blks_vacuumed,  
index_vacuum_count, max_dead_tuples, num_dead_tuples  
FROM pg_stat_progress_vacuum WHERE datname = current_database();
```

*Record 1*

now	2019-08-01 11:32:15.300526-04
pid	4204
table	esquema.tabela
phase	scanning heap
heap_blks_total	134007
heap_blks_scanned	105442
heap_blks_vacuumed	0
index_vacuum_count	0
max_dead_tuples	291
num_dead_tuples	0

# Reporting vacuum progress (4)

```
alvherre=# \t
alvherre=# \pset tuples_only on
alvherre=# SELECT .. FROM pg_stat_progress_vacuum \watch 0,1
```

# VACUUM: operation phases

1. initializing
2. scanning heap
3. vacuuming indexes
4. vacuuming heap
5. cleaning up indexes
6. truncating heap
7. performing final cleanup

# VACUUM: operation phases

1. initializing
2. scanning heap  
→ `dead_tuples[0 ... maintenance_work_mem]`
3. vacuuming indexes
4. vacuuming heap
5. cleaning up indexes
6. truncating heap
7. performing final cleanup

<i>Record 1</i>	
now	...
pid	...
table	...
phase	...
heap_blks_total	134007
heap_blks_scanned	105442
heap_blks_vacuumed	0
index_vacuum_count	0
max_dead_tuples	291
num_dead_tuples	0

# VACUUM: operation phases

1. initializing
2. scanning heap
3. vacuuming indexes
4. vacuuming heap
5. cleaning up indexes
6. truncating heap
7. performing final cleanup  
→ FSM update

# VACUUM: operation phases

1. initializing
2. scanning heap
3. vacuuming indexes
4. vacuuming heap
5. cleaning up indexes
6. truncating heap
  - requires access exclusive lock
  - step not done if unavailable
7. performing final cleanup
  - FSM update

# Reporting CREATE INDEX / REINDEX progress

- PostgreSQL 12
- Report progress of CREATE INDEX operations
  - <http://git.postgresql.org/pg/commitdiff/ab0dfc961b6a>
  - Álvaro Herrera
- Report progress of REINDEX operations
  - <http://git.postgresql.org/pg/commitdiff/03f9e5cba0ee>
  - Peter Eisentraut



# Report of CREATE INDEX

```
SELECT ... FROM pg_stat_progress_create_index ... \watch 1
```

pid	1209
relid	esquema.tabela
index_relid	35684
command	CREATE INDEX CONCURRENTLY
phase	building index: scanning table
lockers_total	0
lockers_done	0
current_locker_pid	0
blocks_total	44248
blocks_done	17627
tuples_total	0
tuples_done	0
partitions_total	0
partitions_done	0

# Operation phases of CREATE INDEX / REINDEX

1. initializing
2. waiting for writers before build
3. building index
4. waiting for writers before validation
5. index validation: scanning index
6. index validation: sorting tuples
7. index validation: scanning table
8. waiting for old snapshots
9. waiting for readers before marking dead
10. waiting for readers before dropping

# Operation phases of CREATE INDEX / REINDEX

1. initializing
2. waiting for writers before build
3. building index
4. waiting for writers before validation
5. index validation: scanning index
6. index validation: sorting tuples
7. index validation: scanning table
8. waiting for old snapshots
9. waiting for readers before marking dead
10. waiting for readers before dropping

phase	....
lockers_total	0
lockers_done	0
current_locker_pid	0
blocks_total	44248
blocks_done	17627
tuples_total	0
tuples_done	0
partitions_total	0
partitions_done	0

# Build phases for B-Tree indexes

1. initializing
2. scanning table
3. sorting live tuples
4. sorting dead tuples
5. loading tuples in tree

# Reporting CLUSTER / VACUUM FULL progress

- PostgreSQL 12
- Add progress reporting for CLUSTER and VACUUM FULL.
  - <http://git.postgresql.org/pg/commitdiff/6f97457e0ddd>
  - Tatsuro Yamada

# Reporting cluster progress

```
alvherre=# SELECT * FROM pg_stat_progress_cluster \watch 1
```

pid	1209
table	esquema.tabela
command	VACUUM FULL
phase	seq scanning heap
cluster_index_relid	0
heap_tuples_scanned	8064358
heap_tuples_written	8064358
heap_blks_total	44248
heap_blks_scanned	35684
index_rebuild_count	0

# CLUSTER: operation phases

1. initializing
2. seq scanning heap
3. index scanning heap
4. sorting tuples
5. writing new heap
6. swapping relation files
7. rebuilding index
8. performing final cleanup

# Reporting ANALYZE progress

- Patch submitted for PostgreSQL 13
- <https://postgr.es/m/20190621185207.GA27929@alvherre.pgsql>



# Questions?

Thanks for listening!

# Appendix: Implementation

- Set-returning function `pg_stat_get_progress_info(text)`
- Returns raw metrics
- View definitions (`pg_stat_progress_vacuum` etc) transform metrics into user-readable parameters
- PostgreSQL C code injects metrics into `pgstat` system

```
pgstat_progress_start_command(command type);  
pgstat_progress_update_param(4, 158);  
pgstat_progress_update_param(PROGRESS_ANALYZE_PHASE,  
                             PROGRESS_ANALYZE_PHASE_SCAN_TABLE);  
pgstat_progress_end_command();
```