

# The MERGE Command

Álvaro Herrera  
PostgreSQL developer



<https://www.EnterpriseDB.com/>

Prague PostgreSQL Developer Day  
February 2023

# So what is MERGE anyway?

- A “new” SQL DML command

ISO/IEC 9075-2:2016(E)

**14.12** <merge statement>

## **14.12** <merge statement>

*This Subclause is modified by Subclause 14.5, “<merge statement>”, in ISO/IEC 9075-14.*

### **Function**

Conditionally update and/or delete rows of a table and/or insert new rows into a table.

- Introduced in SQL:2003
- Is it UPSERT?

# So what is MERGE anyway?

- A “new” SQL DML command

ISO/IEC 9075-2:2016(E)

**14.12** <merge statement>

## **14.12** <merge statement>

*This Subclause is modified by Subclause 14.5, “<merge statement>”, in ISO/IEC 9075-14.*

### **Function**

Conditionally update and/or delete rows of a table and/or insert new rows into a table.

- Introduced in SQL:2003
- Is it UPSERT?

# What MERGE is not

- MERGE is not UPSERT
- (PostgreSQL ended up introducing non-standard INSERT ON CONFLICT UPDATE in 9.5 (2016))

```
INSERT INTO wines (winery, brand, variety, year, bottles)
    VALUES ('Concha y Toro', 'Sunrise', 'Chardonnay', 2021, 96)
ON CONFLICT (winery, brand, variety, year) DO
    UPDATE SET bottles = wines.bottles + EXCLUDED.bottles;
```

- 👍 Special (index-)row locking is used
- 👍 Fast and “concurrency correct”
- 🗑 It can't do DELETE
- 🗑 It is not standard and very different from MERGE
- 🗑 Concurrency considerations are different
- )

# What MERGE is not

- MERGE is not UPSERT
- (PostgreSQL ended up introducing non-standard INSERT ON CONFLICT UPDATE in 9.5 (2016))

```
INSERT INTO wines (winery, brand, variety, year, bottles)
    VALUES ('Concha y Toro', 'Sunrise', 'Chardonnay', 2021, 96)
ON CONFLICT (winery, brand, variety, year) DO
    UPDATE SET bottles = wines.bottles + EXCLUDED.bottles;
```

- 👍 Special (index-)row locking is used
- 👍 Fast and “concurrency correct”
- 🗑 It can't do DELETE
- 🗑 It is not standard and very different from MERGE
- 🗑 Concurrency considerations are different

• )

# What MERGE is not

- MERGE is not UPSERT
- (PostgreSQL ended up introducing non-standard INSERT ON CONFLICT UPDATE in 9.5 (2016))

```
INSERT INTO wines (winery, brand, variety, year, bottles)
    VALUES ('Concha y Toro', 'Sunrise', 'Chardonnay', 2021, 96)
ON CONFLICT (winery, brand, variety, year) DO
    UPDATE SET bottles = wines.bottles + EXCLUDED.bottles;
```

- 👍 Special (index-)row locking is used
- 👍 Fast and “concurrency correct”
- 🗨 It can't do DELETE
- 🗨 It is not standard and very different from MERGE
- 🗨 Concurrency considerations are different
- )

# MERGE history

- First implementation attempt in PostgreSQL 11 (2018)
- Fully implemented in PostgreSQL 15 (2022)
- Other RDBMS systems have it
  - Oracle™ 11 had it (2007?)
  - SQL Server™ 2008 had it
- ... so Postgres having it, improves chances of migration
  - (even though there are still some differences)

# MERGE history

- First implementation attempt in PostgreSQL 11 (2018)
- Fully implemented in PostgreSQL 15 (2022)
- Other RDBMS systems have it
  - Oracle™ 11 had it (2007?)
  - SQL Server™ 2008 had it
- ... so Postgres having it, improves chances of migration
  - (even though there are still some differences)



# MERGE example

```
CREATE TABLE wines (winery text, brand text, variety text, year int,  
                    bottles int);  
ALTER TABLE wines ADD UNIQUE (winery, brand, variety, year);  
CREATE TABLE shipment (LIKE wines);  
  
INSERT INTO shipment VALUES  
    ('Concha y Toro', 'Sunrise', 'Chardonnay', 2021, 96),  
    ('Concha y Toro', 'Sunrise', 'Merlot', 2022, 120),  
    ('Concha y Toro', 'Marqués de Casa y Concha', 'Carmenere',  
     2021, 48),  
    ('Concha y Toro', 'Casillero del Diablo', 'Cabernet Sauvignon',  
     2019, 240);
```

# MERGE example

```
CREATE TABLE wines (winery text, brand text, variety text, year int, bottles int);  
CREATE TABLE shipment (LIKE wines);
```

```
MERGE INTO wines AS w  
  USING shipment AS s  
    ON (w.winery, w.brand, w.variety, w.year) =  
       (s.winery, s.brand, s.variety, s.year)
```

```
WHEN MATCHED THEN  
  UPDATE SET bottles = w.bottles + s.bottles
```

```
WHEN NOT MATCHED THEN  
  INSERT (winery, brand, variety, year, bottles)  
  VALUES (s.winery, s.brand, s.variety, s.year, s.bottles);
```

# MERGE syntax

```
[ WITH clause ]
```

```
MERGE INTO target_table  
    USING { source table or query }  
    ON { join condition }  
  
    WHEN MATCHED [ AND expression ] THEN  
        UPDATE SET { columns / values }  
  
    WHEN MATCHED [ AND expression ] THEN  
        DELETE  
  
    WHEN NOT MATCHED [ AND expression ] THEN  
        INSERT { columns / values }  
;
```

# Multiple WHEN-clause example

```
MERGE INTO wines w
  USING shipment s
    ON (w.winery, ...) = (s.winery, ...)

  WHEN MATCHED AND (w.bottles + s.bottles) <= 0
    DELETE
  WHEN MATCHED AND (w.bottles + s.bottles) > 1200
    UPDATE SET on_sale = true, bottles = w.bottles+s.bottles
  WHEN MATCHED AND (w.bottles + s.bottles) < 120
    UPDATE SET on_sale = false, bottles = w.bottles+s.bottles
  WHEN MATCHED
    UPDATE SET bottles = w.bottles + s.bottles
  WHEN NOT MATCHED AND s.bottles < 12 THEN
    DO NOTHING
  WHEN NOT MATCHED THEN
    INSERT ( ... ) VALUES ( ... )
```

# Dealing with concurrency

- Easiest is to not run two MERGEs concurrently targetting the same table
  - (perhaps: LOCK TABLE wines IN SHARE MODE)
- If you must have two, make them not have WHEN NOT MATCHED THEN INSERT clauses
- If you must allow concurrent insertion, rewrite to INSERT ON CONFLICT UPDATE
- If you do not lock the table, test extensively for concurrent scenarios
  - plan to spend ~~at least twice as long~~ much longer testing than developing

# Dealing with concurrency

- Easiest is to not run two MERGEs concurrently targetting the same table
  - (perhaps: LOCK TABLE wines IN SHARE MODE)
- If you must have two, make them not have WHEN NOT MATCHED THEN INSERT clauses
- If you must allow concurrent insertion, rewrite to INSERT ON CONFLICT UPDATE
- If you do not lock the table, test extensively for concurrent scenarios
  - plan to spend ~~at least twice as long~~ much longer testing than developing

## Duplicates in MERGE source

What if my source table has “duplicates”?

```
-- update
```

```
ERROR:  MERGE command cannot affect row a second time
```

```
-- insert
```

```
ERROR:  duplicate key value violates unique constraint "winery"
```

```
DETAIL:  Key (winery, brand, variety, year)=(Concha y Toro,
```

## Duplicates in MERGE source

What if my source table has “duplicates”?

```
-- update
```

```
ERROR:  MERGE command cannot affect row a second time
```

```
-- insert
```

```
ERROR:  duplicate key value violates unique constraint "win
```

```
DETAIL:  Key (winery, brand, variety, year)=(Concha y Toro,
```



## Duplicates in MERGE source (2)

```
MERGE INTO wines AS w
  USING (  SELECT winery, brand, variety, year,
                sum(bottles) as bottles
          FROM shipment
          GROUP BY winery, brand, variety, year
        ) AS s
  ON (w.winery, w.brand, w.variety, w.year) =
     (s.winery, s.brand, s.variety, s.year)
  WHEN MATCHED THEN
    UPDATE SET bottles = w.bottles + s.bottles
  WHEN NOT MATCHED THEN
    INSERT (winery, brand, variety, year, bottles)
    VALUES (s.winery, s.brand, s.variety, s.year, s.bottles)
;
```

## MERGE example 3

What if I run MERGE twice with the same source?

```
ALTER TABLE shipment ADD COLUMN marked timestamp with time zone;
```

```
WITH unmarked_shipment AS  
  (UPDATE shipment SET marked = now() WHERE marked IS NULL  
   RETURNING winery, brand, variety, year, bottles)
```

```
MERGE INTO wines AS w  
  USING unmarked_shipment AS s  
    ON (w.winery, w.brand, w.variety, w.year) =  
       (s.winery, s.brand, s.variety, s.year)
```

```
WHEN MATCHED THEN  
  UPDATE SET bottles = w.bottles + s.bottles
```

```
WHEN NOT MATCHED THEN  
  INSERT (winery, brand, variety, year, bottles)  
  VALUES (s.winery, s.brand, s.variety, s.year, s.bottles);
```

## MERGE example 3

What if I run MERGE twice with the same source?

```
ALTER TABLE shipment ADD COLUMN marked timestamp with time zone;
```

```
WITH unmarked_shipment AS  
  (UPDATE shipment SET marked = now() WHERE marked IS NULL  
   RETURNING winery, brand, variety, year, bottles)
```

```
MERGE INTO wines AS w  
  USING unmarked_shipment AS s  
    ON (w.winery, w.brand, w.variety, w.year) =  
       (s.winery, s.brand, s.variety, s.year)
```

```
WHEN MATCHED THEN  
  UPDATE SET bottles = w.bottles + s.bottles
```

```
WHEN NOT MATCHED THEN  
  INSERT (winery, brand, variety, year, bottles)  
  VALUES (s.winery, s.brand, s.variety, s.year, s.bottles);
```

## MERGE example 3

What if I run MERGE twice with the same source?

```
ALTER TABLE shipment ADD COLUMN marked timestamp with time zone;
```

```
WITH unmarked_shipment AS
```

```
  (UPDATE shipment SET marked = now() WHERE marked IS NULL  
   RETURNING winery, brand, variety, year, bottles)
```

```
MERGE INTO wines AS w
```

```
  USING unmarked_shipment AS s
```

```
    ON (w.winery, w.brand, w.variety, w.year) =  
       (s.winery, s.brand, s.variety, s.year)
```

```
WHEN MATCHED THEN
```

```
  UPDATE SET bottles = w.bottles + s.bottles
```

```
WHEN NOT MATCHED THEN
```

```
  INSERT (winery, brand, variety, year, bottles)
```

```
  VALUES (s.winery, s.brand, s.variety, s.year, s.bottles);
```

# Combining both examples

```
WITH unmarked_shipment AS
  (UPDATE shipment SET marked = now() WHERE marked IS NULL
   RETURNING winery, brand, variety, year, bottles)
MERGE INTO wines AS w
  USING (SELECT winery, brand, variety, year,
               sum(bottles) as bottles
        FROM unmarked_shipment
        GROUP BY winery, brand, variety, year) AS s
  ON (w.winery, w.brand, w.variety, w.year) =
     (s.winery, s.brand, s.variety, s.year)
WHEN MATCHED THEN
  UPDATE SET bottles = w.bottles + s.bottles
WHEN NOT MATCHED THEN
  INSERT (winery, brand, variety, year, bottles)
  VALUES (s.winery, s.brand, s.variety, s.year, s.bottles)
;
```

# Trigger behavior

- Boring (it behaves as you'd expect)
  - BEFORE EACH STATEMENT
    - INSERT
    - UPDATE
    - DELETE
  - BEFORE EACH ROW
    - Each row as scanned by the join
  - AFTER EACH ROW
    - Each row in the same order as above
  - AFTER EACH STATEMENT
    - DELETE
    - UPDATE
    - INSERT

## FOR STATEMENT triggers – transition table example

```
CREATE TABLE wine_audit (op varchar(1), datetime timestamptz,  
                           oldrow jsonb, newrow jsonb);  
  
CREATE FUNCTION wine_audit() RETURNS trigger LANGUAGE plpgsql AS $$  
BEGIN  
    IF (TG_OP = 'DELETE') THEN  
        INSERT INTO wine_audit  
            SELECT 'D', now(), row_to_json(o), NULL FROM old_table o;  
    ELSIF (TG_OP = 'INSERT') THEN  
        INSERT INTO wine_audit  
            SELECT 'I', now(), NULL, row_to_json(n) FROM new_table n;  
    ELSIF (TG_OP = 'UPDATE') THEN  
        ... -- for later  
    END IF;  
    RETURN NULL;  
END;  
$$;
```

# Transition table – the triggers

```
CREATE TRIGGER wine_update
  AFTER UPDATE ON wines
  REFERENCING OLD TABLE AS old_table NEW TABLE AS new_table
  FOR EACH STATEMENT EXECUTE FUNCTION wine_audit();
```

```
CREATE TRIGGER wine_insert
  AFTER INSERT ON wines
  REFERENCING NEW TABLE AS new_table
  FOR EACH STATEMENT EXECUTE FUNCTION wine_audit();
```

```
CREATE TRIGGER wine_delete
  AFTER DELETE ON wines
  REFERENCING OLD TABLE AS old_table
  FOR EACH STATEMENT EXECUTE FUNCTION wine_audit();
```



## Transition table – UPDATE part

```
DECLARE
    oldrec record;
    newrec jsonb;
    i      integer := 0;
BEGIN
    FOR oldrec IN SELECT * FROM old_table LOOP
        newrec := row_to_json(n) FROM new_table n OFFSET i LIMIT 1;
        i := i + 1;
        INSERT INTO wine_audit
            SELECT 'U', now(), row_to_json(oldrec), newrec;
    END LOOP;
END;
```

# Transition table – complete example

```
CREATE FUNCTION wine_audit() RETURNS trigger LANGUAGE plpgsql AS $$
BEGIN
  IF (TG_OP = 'DELETE') THEN
    INSERT INTO wine_audit
      SELECT 'D', now(), row_to_json(o), NULL FROM old_table o;
  ELSEIF (TG_OP = 'INSERT') THEN
    INSERT INTO wine_audit
      SELECT 'I', now(), NULL, row_to_json(n) FROM new_table n;
  ELSEIF (TG_OP = 'UPDATE') THEN
    DECLARE
      oldrec record;
      newrec jsonb;
      i integer := 0;
    BEGIN
      FOR oldrec IN SELECT * FROM old_table LOOP
        newrec := row_to_json(n) FROM new_table n OFFSET i LIMIT 1;
        i := i + 1;
        INSERT INTO wine_audit
          SELECT 'U', now(), row_to_json(oldrec), newrec;
      END LOOP;
    END;
  END IF;
  RETURN NULL;
END;
$$;
```

# Permissions

- Permissions are required according to operations specified
- ... even if at run-time the operations are not executed
- Requirements are like normal INSERT/UPDATE/DELETE

# Possible target types

- Supported
  - Regular tables
  - Partitioned tables
  - Tables with “legacy” inheritance
- Not yet supported
  - Foreign tables
  - Materialized views
  - Updatable views

# Possible target types

- Supported
  - Regular tables
  - Partitioned tables
  - Tables with “legacy” inheritance
- Not yet supported
  - Foreign tables
  - Materialized views
  - Updatable views

# Development Credits

- Simon Riggs
- Pavan Deolasee
- Amit Langote
- Álvaro Herrera

# Future projects

All by Dean Rasheed:

- Support for updatable views
- Support for RETURNING
- Support for WHEN NOT MATCHED BY SOURCE

## Questions?

Álvaro Herrera, EDB

`alvherre@alvh.no-ip.org`

`https://alvherre.cl/`

Mastodon: `https://lile.cl/@alvherre/`